

CHAPTER 7 Equijoins

Oracle SQL By Example, Fourth Edition by Alice Rischert. Published by Prentice Hall. Copyright © 2008 by Pearson Education, Inc.

CHAPTER OBJECTIVES

In this chapter, you will learn about:

- ▶ The Two-Table Join
- ▶ Joining Three or More Tables

So far, you have written SQL statements against a single table. In this chapter, you will learn about joining tables, one of the most important aspects of the SQL language. The *equijoin*, which is by far the most common form of join, allows you to connect two or more tables. Equijoins are based on equality of values in one or more columns. You will learn about other types of joins in [Chapter 10](#), “Complex Joins.”

285

286

Lab 7.1 The Two-Table Join

LAB OBJECTIVES

After this lab, you will be able to:

- ▶ Write Simple Join Constructs
- ▶ Narrow Down Your Result Set
- ▶ Understand the Cartesian Product

In this lab, you will join information from two tables into one meaningful result. Suppose you want to list the course number, course description, section number, location, and instructor ID for each section. This data is found in two separate tables: The course number and description are in the COURSE table; the SECTION table contains the course number, section number, location, and instructor ID. One approach is to query the individual tables and record the results on paper, then match every course number in the COURSE table with the corresponding course number in the SECTION table. The other approach is to formulate a SQL statement that accomplishes the join for you.

[Figure 7.1](#) shows part of the COURSE table. Missing columns and rows are indicated with an ellipsis (...). The primary key of the COURSE table is COURSE_NO.

FIGURE 7.1 Excerpt of the COURSE table

COURSE_NO	DESCRIPTION						MODIFIED_DATE
10	Technology Concepts	05-APR-07
20	Intro to Information Systems	05-APR-07
25	Intro to Programming	05-APR-07
80	Programming Techniques	05-APR-07
100	Hands-On Windows	05-APR-07
120	Intro to Java Programming	05-APR-07
122	Intermediate Java Programming	05-APR-07
124	Advanced Java Programming	05-APR-07
125	Java Developer I	05-APR-07
130	Intro to Unix	05-APR-07
132	Design of Unix Admin	05-APR-07

286

287

[Figure 7.2](#) shows part of the SECTION table. The COURSE_NO column is the foreign key to the COURSE table.

FIGURE 7.2 Excerpt of the SECTION table

SECTION_ID	COURSE_NO	SECTION_NO	LOCATION
80	10	2	L214
81	20	2	L210
82	20	4	L214
83	20	7	L509
84	20	8	L210
85	25	1	M311
86	25	2	L210
87	25	3	L507
88	25	4	L214
89	25	5	L509
90	25	6	L509
91	25	7	L210
92	25	8	L509
93	25	9	L507
141	100	1	L214
142	100	2	L500

Examine the result set listed in [Figure 7.3](#). For example, for course number 10, one section exists in the SECTION table. The result of the match is one row. Course number 20, Intro to Information Systems, has multiple rows in the section table because there are multiple classes/sections for the same course. Course number 80, Programming Techniques, is missing from the result. This course number has no matching entry in the SECTION table, and therefore this row is not in the result.

FIGURE 7.3 Result of the join between the COURSE and SECTION tables

COURSE_NO	SECTION_NO	DESCRIPTION	LOCATION	INSTRUCTOR_ID
10	2	Technology Concepts	L214	102
20	2	Intro to Information Systems	L210	103
20	4	Intro to Information Systems	L214	104
20	7	Intro to Information Systems	L509	105
20	8	Intro to Information Systems	L210	106
25	1	Intro to Programming	M311	107
25	2	Intro to Programming	L210	108
25	3	Intro to Programming	L507	101
25	4	Intro to Programming	L214	102
25	5	Intro to Programming	L509	103
25	6	Intro to Programming	L509	104
25	7	Intro to Programming	L210	105
25	8	Intro to Programming	L509	106
25	9	Intro to Programming	L507	107
100	1	Hands-On Windows	L214	102
100	2	Hands-On Windows	L500	103
100	3	Hands-On Windows	L509	104

Steps to Formulate the SQL Statement

Before you write the SQL join statement, first choose the columns you want to include in the result. Next, determine the tables to which the columns belong. Then identify the common columns between the tables.

Finally, determine whether there is a one-to-one or a one-to-many relationship among the column values. Joins are typically used to join between the primary key and the foreign key. In the previous example, the COURSE_NO column in the COURSE table is the primary key, and the column COURSE_NO in the SECTION table is the foreign key. This represents a one-to-many relationship between the tables. (Joining tables related through a many-to-many relationship yields a *Cartesian product*. You'll learn more about the Cartesian product later in this chapter.)

Following is the SQL statement that achieves the result shown in [Figure 7.3](#). It looks much like the SELECT statements you have written previously, but two tables, separated by commas, are listed in the FROM clause.

```
SELECT course. course_no,  
       section_no, description,  
       location, instructor_id  
FROM course, section
```

```
WHERE course. course_no =  
      section. course_no
```

The WHERE clause formulates the *join criteria*, also called the *join condition*, between the two tables using the common COURSE_NO column. Because this is an equijoin, the values in the common columns must equal each other for a row to be displayed in the result set. Each COURSE_NO value from the COURSE table must match a COURSE_NO value from the SECTION table. To differentiate between columns of the same name, *qualify* the columns by prefixing the column with the table name and a period. Otherwise, Oracle returns the error ORA-00918: column ambiguously defined.

Instead of displaying the COURSE_NO column from the COURSE table in the SELECT list, you can use the COURSE_NO column from the SECTION table. Because it is an equijoin, it returns the same result.



The order in which the tables are listed in the FROM has no effect on the query result. But the join order can become important when optimizing SQL statement, and

you will learn about this in [Chapter 18](#), “SQL Optimization.”

Table Aliases

Instead of using the table name as a prefix to differentiate between the columns, you can use a *table alias*, which qualifies the table using a short abbreviation.

```
SELECT c.course_no, s.section_no,  
       c.description,  
       s.location, s.instructor_id  
FROM   course c, section s  
WHERE  c.course_no = s.course_no
```

288

289

The table alias names are arbitrary. However, you cannot use any Oracle *reserved words*. (Reserved words have a special meaning in SQL or in the Oracle database and are typically associated with a SQL command. For example, SELECT and WHERE are reserved words.) It is best to keep the name short and simple, as in this example, where the COURSE table has the alias c, and the SECTION table has the alias s.

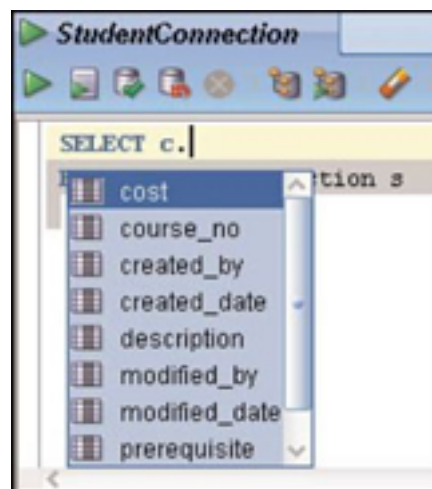


To easily identify the source table of a column and to improve the readability of a join statement, it is best to qualify all column names with the table alias.

Furthermore, this avoids any future ambiguities that may arise if a new column with the same name is added later. Without a qualified table alias, a subsequently issued SQL statement referencing both tables results in the Oracle error ORA-00918: column ambiguously defined.

When you define a table alias in SQL Developer, any subsequent reference of the alias brings up a list of columns for the table, helping you remember the column names and avoid column misspellings (see [Figure 7.4](#)). Refer to [Chapter 2](#), “SQL: The Basics,” for more information regarding the code completion feature.

FIGURE 7.4 Column list of the COURSE table



Narrowing Down Your Result Set

The previous SQL statement lists all the rows in the SECTION and COURSE tables with matching COURSE_NO values. If you want to narrow down the criteria to specific rows, you can expand the WHERE clause to include additional conditions. The next statement chooses only those courses and their respective sections where the DESCRIPTION column starts with the text Intro to.

```
SELECT c.course_no, s.section_no,  
       c.description,  
       s.location, s.instructor_id  
FROM   course c, section s  
WHERE  c.course_no = s.course_no  
       AND c.description LIKE 'Intro to  
       %'
```

289

290

Nulls and Joins

In an equijoin, a null value in the common column has the effect of not including the row in the result. Look at the foreign key column ZIP on the INSTRUCTOR table, which allows nulls.

First, query for records with a null value.

```
SELECT instructor_id, zip, last_name, first_name
FROM instructor
WHERE zip IS NULL
INSTRUCTOR_ID ZIP    LAST_NAME  FIRST_NAME
-----
          110      Willig      Irene

1 row selected.
```

Next, formulate the join to the ZIPCODE table via the ZIP column. Observe that instructor Irene Willig does not appear in the result.

```
SELECT i.instructor_id, i.zip, i.last_name, i.first_name
FROM instructor i, zipcode z
WHERE i.zip = z.zip
INSTRUCTOR_ID ZIP    LAST_NAME  FIRST_NAME
-----
          101 10015 Hanks      Fernand
          105 10015 Morris     Anita
          109 10015 Chow       Rick
          102 10025 Wojick     Tom
          103 10025 Schorin    Nina
          106 10025 Smythe     Todd
          108 10025 Lowry      Charles
          107 10005 Frantzen   Marilyn
          104 10035 Pertez     Gary

9 rows selected.
```

A null value is not equal to any other value, including another null value. In this case, the zip code of Irene Willig's record is null; therefore, this row is not included

in the result. In [Chapter 10](#), you will learn how to include null values by formulating an *outer join* condition.

ANSI Join Syntax

Starting with Oracle 9i, Oracle implemented a number of additions to SQL to conform to many aspects of the ANSI SQL/92 and SQL:1999 standards. The advantage of the ANSI join syntax over the traditional comma-separated tables FROM clause is that SQL queries can run unmodified against other non-Oracle, ANSI-compliant databases.

290

291

INNER JOINS

The term *inner join* is used to express a join that satisfies the join condition. Typically, the join condition is based on equality, thus creating an equijoin. (The inner join is in contrast to the outer join. Besides the matched rows, the outer join also includes the unmatched rows from two tables.)

The ANSI syntax has a number of differences from the previously discussed join syntax. For example, the JOIN keyword replaces the comma between the tables and identifies the tables to be joined.

The keyword **INNER** is optional and typically omitted. To express a join condition, you can specify either the **USING** condition or the **ON** condition.

THE USING CONDITION

The **USING** condition, also referred as the **USING** clause, identifies the common column between the tables. Here the common column is the **COURSE_NO** column, which has the same name and compatible data type in both tables. An equijoin is always assumed with the **USING** clause.

```
SELECT course_no, s.section_no,  
       c.description,  
       s.location, s.instructor_id  
FROM   course c JOIN section s  
       USING (course_no)
```

Alternatively, you can include the optional **INNER** keyword, but as mentioned, this keyword is usually omitted.

```
SELECT course_no, s.section_no,  
       c.description,  
       s.location, s.instructor_id  
FROM   course c INNER JOIN section  
       s
```

```
USING (course_no)
```

The following query does not execute because you cannot use a table alias name with this syntax. The USING syntax implies that the column names are identical. The Oracle error identifies the C.COURSE_NO column as the column in the USING clause that has the problem.

```
SQL> SELECT course_no,  
s.section_no, c.description,  
2 s.location, s.instructor_id  
3 FROM course c JOIN section s  
4 USING (c.course_no)  
5 /  
USING (c.course_no)  
*
```

ERROR at line 4:

**ORA-01748: only simple column
names allowed here**

The next query shows the COURSE_NO column in the SELECT list prefixed with the alias name, thus resulting in an error as well. This alias must also be eliminated to successfully run the query.

```
SQL> SELECT c.course_no,  
s.section_no, c.description,
```

291

292

```
2 s.location, s.instructor_id
3 FROM course c JOIN section s
4* USING (course_no)
SQL> /
SELECT c.course_no,
s.section_no, c.description,
*
ERROR at line 1:
ORA-25154: column part of USING
clause cannot have qualifier
```

The next example illustrates the error you receive when you omit the parentheses around the column COURSE_NO in the USING clause.

```
SQL> SELECT c.course_no,
s.section_no, c.description,
2 s.location, s.instructor_id
3 FROM course c JOIN section s
4 USING course_no
5 /
USING course_no
*
ERROR at line 4:
ORA-00906: missing left parenthesis
```

THE ON CONDITION

In case the column names on the tables are different, you use the ON condition, also referred to as the ON clause. The next query is identical in functionality to the previous query, but it is now expressed with the ON condition, and the column name is qualified with the alias both in the SELECT list and in the ON condition. This syntax allows for conditions other than equality and different column names; you will see many such examples in [Chapter 10](#).

```
SELECT c.course_no, s.section_no,  
       c.description,  
       s.location, s.instructor_id  
FROM   course c JOIN section s  
       ON (c.course_no = s.course_no)
```

The pair of parentheses around the ON condition is optional. When comparing this syntax to the traditional join syntax, there are not many differences other than the ON clause and the JOIN keyword.

ADDITIONAL WHERE CLAUSE CONDITIONS

The ON and USING conditions let you specify the join condition separately from any other WHERE condition. One of the advantages of the ANSI join syntax is that it separates the join condition from the filtering WHERE clause condition.

```
SELECT c.course_no, s.section_no,  
       c.description,  
       s.location, s.instructor_id  
FROM   course c JOIN section s  
       ON (c.course_no = s.course_no)  
WHERE description LIKE 'B%'
```

292

293

NATURAL JOINS

A natural join joins tables based on the columns with the same name and data type. Here there is no need to prefix the column name with the table alias, and the join is indicated with the keywords NATURAL JOIN. There is not even a mention of which column(s) to join. This syntax figures out the common columns between the tables. Any use of the ON clause or the USING clause is not allowed with the NATURAL JOIN keywords, and the common columns cannot list a table alias.

```
SELECT course_no, s.section_no,  
c.description,  
s.location, s.instructor_id  
FROM course c NATURAL JOIN section  
S  
no rows selected
```

You might be surprised that the query does not return any result. However, the COURSE_NO column is not the only column with a common name. The columns CREATED_BY, CREATED_DATE, MODIFIED_BY, and MODIFIED_DATE are also common to both tables. These columns record the name of the last user updating a row and the original user creating the row, including the respective date and time. The SQL statement does not return any results because there are no rows that have identical values for all five common columns.



Using the natural join within a program is somewhat risky. There is always a chance for adding to the table columns that happen to have the same column name, and then you may not get the desired result. Therefore,

the natural join works best for ad hoc queries but not for repeated use within programs.

The Cartesian Product

The Cartesian product is rarely useful in the real world. It usually indicates either that the WHERE clause has no joining columns or that multiple rows from one table match multiple rows in another table; in other words, it indicates a many-to-many relationship.

To illustrate the multiplication effect of a Cartesian product, the following query joins the INSTRUCTOR table with the SECTION table. The INSTRUCTOR table contains 10 rows; the SECTION table has 78 rows. The multiplication of all the possible combinations results in 780 rows.

```
SELECT COUNT(*)
FROM section, instructor
COUNT(*)
-----
       780

1 row selected.
```

293

294

Following is a partial listing of the rows, showing all the different combinations of values between the two tables.

```
SELECT s.instructor_id s_instructor_id,
       i.instructor_id i_instructor_id
FROM section s, instructor i
S_INSTRUCTOR_ID I_INSTRUCTOR_ID
-----
          101          101
          101          101
          101          101
          101          101
          101          101
          101          101
          101          101
          101          101
          101          101
          101          101
          101          102
          101          102
          101          102
...
          108          110
          101          110
```

780 rows selected.

In SQL*Plus, if you want to stop and examine the rows one screen at a time, you can use the SQL*Plus SET PAUSE ON command. [Appendix C](#), “SQL*Plus Command Reference,” provides more details about the various SQL*Plus commands.

THE ANSI STANDARD CROSS-JOIN

To formulate a Cartesian product using the ANSI JOIN syntax, you use the keyword **CROSS JOIN** in place of the comma between the two tables. Because of the nature of a cross-join as a combination of all possible values, the SQL statement does not have join criteria. The result is obviously identical to that of the Cartesian product.

```
SELECT COUNT (*)
FROM section CROSS JOIN instructor
```

LAB 7.1 EXERCISES

- a) For all students, display last name, city, state, and zip code. Show the result ordered by zip code.
- b) Select the first and last names of all enrolled students and order by last name in ascending order.
- c) Execute the following SQL statement. Explain your observations about the WHERE clause and the resulting output.

```
SELECT c.course_no,
       c.description, s.section_no
```

294
295

```
FROM course c, section s
WHERE c.course_no =
s.course_no
AND c.prerequisite IS NULL
ORDER BY c.course_no,
s.section_no
```

- d) Select the student ID, course number, enrollment date, and section ID for students who enrolled in course number 20 on January 30, 2007.
- e) Select the students and instructors who live in the same zip code by joining on the common ZIP column. Order the result by the STUDENT_ID and INSTRUCTOR_ID columns. What do you observe?

LAB 7.1 EXERCISE ANSWERS

- a) For all students, display last name, city, state, and zip code. Show the result ordered by zip code.

ANSWER: The common column between the ZIPCODE table and the STUDENT table is the ZIP column. The ZIP column in both tables is defined as NOT NULL. For each row in the

ZIPCODE table there may be zero, one, or multiple students living in one particular zip code. For each student's zip code there must be one matching row in the ZIPCODE table. Only records that satisfy the equality condition of the join are returned.

```
SELECT s.last_name, s.zip, z.state, z.city
FROM student s, zipcode z
WHERE s.zip = z.zip
ORDER BY s.zip
```

LAST_NAME	ZIP	ST	CITY
Norman	01247	MA	North Adams
Kocka	02124	MA	Dorchester
...			
Gilloon	43224	OH	Columbus
Snow	48104	MI	Ann Arbor

268 rows selected.

Because the ZIP column has the same name in both tables, you must qualify the column when you use the traditional join syntax. For simplicity, it is best to use an alias instead of the full table name because it saves you a lot of typing and improves readability. The ORDER BY clause lists the S.ZIP column, as does the SELECT clause. Choosing the Z.ZIP column instead of S.ZIP in the SELECT list or ORDER

BY clause displays the same rows because the values in the two columns have to be equal to be included in the result.

You can also write the query with the ANSI join syntax instead and use the ON condition or the USING condition. If you use the ON condition, you must prefix the common columns with either their full table name or their table alias, if one is specified.

```
SELECT s.last_name, s.zip,
       z.state, z.city
FROM student s JOIN zipcode z
ON (s.zip = z.zip)
ORDER BY s.zip
```

295

If you choose the USING condition instead, do not alias the common column because doing so will cause an error.

```
SELECT s.last_name, zip,
       z.state, z.city
FROM student s JOIN zipcode z
USING (zip)
ORDER BY zip
```

296

- b) Select the first and last names of all enrolled students and order by last name in ascending order.

ANSWER: You need to join the ENROLLMENT and STUDENT tables. Only students who are enrolled have one or multiple rows in the ENROLLMENT table.

```
SELECT s.first_name, s.last_name, s.student_id
FROM student s, enrollment e
WHERE s.student_id = e.student_id
ORDER BY s.last_name
```

FIRST_NAME	LAST_NAME	STUDENT_ID
-----	-----	-----
Mardig	Abdou	119
Suzanne M.	Abid	257
...		
Salewa	Zuckerberg	184
Salewa	Zuckerberg	184
Salewa	Zuckerberg	184
Freedon	annunziato	206

226 rows selected.

Student Salewa Zuckerberg with STUDENT_ID 184 is returned three times. This is because Salewa Zuckerberg is enrolled in three sections. When the SECTION_ID column is included in the SELECT list, this fact becomes evident in the result set.

However, if you are not interested in the `SECTION_ID` and you want to list the names without the duplication, use `DISTINCT` in the `SELECT` statement.

```
SELECT DISTINCT s.first_name,  
s.last_name, s.student_id  
FROM student s, enrollment e  
WHERE s.student_id =  
e.student_id  
ORDER BY s.last_name
```

The `STUDENT_ID` column is required in the `SELECT` clause because there may be students with the same first and last names but who are, in fact, different individuals. The `STUDENT_ID` column differentiates between these students; after all, it's the primary key that is unique to each individual row in the `STUDENT` table.

The student with the last name `annunziato` is the last row. Because the last name is in lowercase, it has a higher sort order. (See Lab 4.3 regarding the sort order values and the `ASCII` function.)

If you use the ANSI syntax, your SQL statement may look similar to the following statement.

```
SELECT s.first_name,  
s.last_name, s.student_id  
FROM student s JOIN enrollment  
e  
ON (s.student_id =  
e.student_id)  
ORDER BY s.last_name
```

296

Or you may write the statement with the USING clause. In this particular query, all aliases are omitted, which has the disadvantage that you cannot easily recognize the source table for each column.

297

```
SELECT first_name, last_name,  
student_id  
FROM student JOIN enrollment  
USING (student_id)  
ORDER BY last_name
```

- c)** Execute the following SQL statement. Explain your observations about the WHERE clause and the resulting output.

```
SELECT c.course_no,  
c.description, s.section_no  
FROM course c, section s
```

```
WHERE c.course_no =
s.course_no
AND c.prerequisite IS NULL
ORDER BY c.course_no,
section_no
```

ANSWER: This query includes both a join condition and a condition that restricts the rows to courses that have no prerequisite. The result is ordered by the course number and the section number.

```
COURSE_NO DESCRIPTION                SECTION_NO
-----
      10 Technology Concepts                2
      20 Intro to Information Systems        2
      ...
     146 Java for C/C++ Programmers         2
     310 Operating Systems                  1

8 rows selected.
```

The COURSE and SECTION tables are joined to obtain the SECTION_NO column. The join requires the equality of values for the COURSE_NO columns in both tables. The courses without a prerequisite are determined with the IS NULL operator.

If the query is written with the ANSI join syntax and the ON clause, you see one advantage of the

ANSI join syntax over the traditional join syntax: The ANSI join distinguishes the join condition from the filtering criteria.

```
SELECT c.course_no,  
       c.description, s.section_no  
FROM   course c JOIN section s  
       ON (c.course_no = s.course_no)  
WHERE  c.prerequisite IS NULL  
ORDER BY c.course_no,  
         section_no
```

- d) Select the student ID, course number, enrollment date, and section ID for students who enrolled in course number 20 on January 30, 2007.

ANSWER: The SECTION and ENROLLMENT tables are joined through their common column: SECTION_ID. This column is the primary key in the SECTION table and the foreign key column in the ENROLLMENT table. The rows are restricted to those records that have a course number of 20 and an enrollment date of January 30, 2007, by including this condition in the WHERE clause.

297

```

SELECT e.student_id, s.course_no,
       TO_CHAR(e.enroll_date, 'MM/DD/YYYY HH:MI PM'),
       e.section_id
FROM enrollment e JOIN section s
  ON (e.section_id = s.section_id)
WHERE s.course_no = 20
      AND e.enroll_date >= TO_DATE('01/30/2007', 'MM/DD/YYYY')
      AND e.enroll_date < TO_DATE('01/31/2007', 'MM/DD/YYYY')
STUDENT_ID COURSE_NO TO_CHAR(ENROLL_DATE SECTION_ID
-----
          103          20 01/30/2007 10:18 AM          81
          104          20 01/30/2007 10:18 AM          81

2 rows selected.

```

Alternatively, you can use the USING clause or the more traditional join syntax, listed here.

```

SELECT e.student_id,
       s.course_no,
       TO_CHAR(e.enroll_date, 'MM/DD/
YYYY HH:MI PM'),
       e.section_id
FROM enrollment e, section s
WHERE e.section_id =
s.section_id
      AND s.course_no = 20
      AND e.enroll_date >=
TO_DATE('01/30/2007', 'MM/DD/
YYYY')
      AND e.enroll_date <
TO_DATE('01/31/2007', 'MM/DD/
YYYY')

```


The WHERE clause considers the date and time values of the ENROLL_DATE column. There are alternative WHERE clause solutions, such as applying the TRUNC function on the ENROLL_DATE column. Refer to [Chapter 5](#), “Date and Conversion Functions,” for many examples of querying and displaying DATE data type columns.

- e) Select the students and instructors who live in the same zip code by joining on the common ZIP column. Order the result by the STUDENT_ID and INSTRUCTOR_ID columns. What do you observe?

ANSWER: When you join the STUDENT and INSTRUCTOR tables, there is a many-to-many relationship, which causes a Cartesian product as a result.

```
SELECT s.student_id, i.instructor_id,
       s.zip, i.zip
FROM student s, instructor i
WHERE s.zip = i.zip
ORDER BY s.student_id, i.instructor_id
```

STUDENT_ID	INSTRUCTOR_ID	ZIP	ZIP
163	102	10025	10025
163	103	10025	10025
163	106	10025	10025
163	108	10025	10025
223	102	10025	10025
223	103	10025	10025
223	106	10025	10025
223	108	10025	10025
399	102	10025	10025
399	103	10025	10025
399	106	10025	10025
399	108	10025	10025

12 rows selected.

Initially, this query and its corresponding result may not strike you as a Cartesian product because the WHERE clause contains a join criteria. However, the relationship between the STUDENT table and the INSTRUCTOR table does not follow the primary key/foreign key path, and therefore a Cartesian product is possible. A look at the schema diagram reveals that no primary key/foreign key relationship exists between the two tables. To further illustrate the many-to-many relationship between the ZIP columns, select the students and instructors living in zip code 10025 in separate SQL statements.

```
SELECT student_id, zip
FROM student
WHERE zip = '10025'
STUDENT_ID ZIP
-----
223 10025
163 10025
399 10025
```

3 rows selected.

```
SELECT instructor_id, zip
FROM instructor
WHERE zip = '10025'
INSTRUCTOR_ID ZIP
-----
102 10025
103 10025
106 10025
108 10025
```

4 rows selected.

These results validate the solution's output: the Cartesian product shows the three student rows multiplied by the four instructors, which results in 12 possible combinations. You can rewrite the query to include the `DISTINCT` keyword to select only the distinct student IDs. You can also write the query with a subquery construct, which avoids the Cartesian product. You will learn about this in [Chapter 8](#), "Subqueries."

JOINING ALONG THE PRIMARY/ FOREIGN KEY PATH

You can join along the primary/foreign key path by joining the `STUDENT` table to the `ENROLLMENT` table, then to the `SECTION` table, and finally to the `INSTRUCTOR` table. This involves a multitable join, discussed in [Lab 7.2](#). However, the result is different from the Cartesian product result because it shows only instructors who teach a section in which the student is enrolled. In other words, an instructor living in zip code 10025 is included in the result only if the instructor teaches that student also living in the same zip code. This is in contrast to the Cartesian product example, which shows all the instructors and students

living in the same zip code, whether the instructor teaches this student or not. You will explore the differences between these two examples once more in the Workshop section at the end of the chapter.

299

300

Lab 7.1 Quiz

In order to test your progress, you should be able to answer the following questions.

1) Find the error(s) in the following SQL statement.

```
1 SELECT stud.last_name, stud.first_name,  
2 stud.zip, zip.zip, zip.state, zip.city,  
3 TO_CHAR(stud.student_id)  
4 FROM student stud, zipcode zip  
5 WHERE stud.student_id = 102  
6 AND zip.zip = '11419'  
7 AND zip.zip = s.zip
```

- a) _____ No error.
- b) _____ This is not an equijoin.
- c) _____ Lines 1, 2, and 3.
- d) _____ Line 4.

e) _____ Lines 5 and 6.

f) _____ Line 7.

2) Find the error(s) in the following SQL statement.

```
1 SELECT s.*, zipcode.zip,  
2 DECODE(s.last_name, 'Smith', szip,  
3 UPPER(s.last_name))  
4 FROM student s, zipcode  
5 WHERE stud.zip = zipcode.zip  
6 AND s.last_name LIKE 'Smi%'
```

a) _____ Lines 1 and 2

b) _____ Lines 1 and 4

c) _____ Line 3

d) _____ Lines 2 and 5

e) _____ Line 4

3) A table alias is the name of a duplicate table stored in memory.

a) _____ True

b) _____ False

4) To equijoin a table with another table involves matching the common column values.

a) _____ True

b) _____ False

5) Find the error(s) in the following SQL statement.

```
1 SELECT TO_CHAR(w.modified_date, 'dd-  
mon-yyyy'),  
2 t.grade_type_code, description,  
3 TO_NUMBER(TO_CHAR(number_per_section))  
4 FROM grade_type t, grade_type_weight w  
5 WHERE t.grade_type_code =  
w.grade_type_code_cd  
6 AND ((t.grade_type_code = 'MT'  
7 OR t.grade_type_code = 'HM'))  
8 AND w.modified_date >=  
9 TO_DATE('01-JAN-2007', 'DD-MON-  
YYYY')
```

a) _____ Lines 1 and 8

300

301

- b) _____Line 4
- c) _____Line 5
- d) _____Lines 6 and 7
- e) _____Lines 5, 6, and 7

6) Given two tables, T1 and T2, and their rows, as shown, which result will be returned by the following query?

```
SELECT t1.val, t2.val, t1.name, t2.location
FROM t1, t2
WHERE t1.val = t2.val
```

Table T1

VAL NAME

A	Jones
B	Smith
C	Zeta
	Miller

Table T2

VAL LOCATION

A	San Diego
B	New York
B	New York
	Phoenix

7) The USING clause of the ANSI join syntax always assumes an equijoin and identical column names.

- a) _____True
- b) _____False

301

8) The NATURAL JOIN keywords and the USING clause of the ANSI join syntax are mutually exclusive.

a) _____ True

b) _____ False

9) The common column used in the join condition must be listed in the SELECT list.

a) _____ True

b) _____ False

ANSWERS APPEAR IN APPENDIX A.

LAB 7. 2 Joining Three or More Tables

LAB OBJECTIVES

After this lab, you will be able to:

- ▶ Join Three or More Tables
- ▶ Join with Multicolumn Join Criteria

You often have to join more than two tables to determine the answer to a query. In this lab, you will practice these

types of joins. In addition, you will join tables with multicolumn keys.

Joining Three or More Tables

The join example at the beginning of the chapter involves two tables: the COURSE and SECTION tables. The following SQL statement repeats that query and the result of the join (see [Figure 7.5](#)). To include the instructor's first and last names, you can expand this statement to join to a third table, the INSTRUCTOR table.

```
SELECT c.course_no, s.section_no,  
       c.description,  
       s.location, s.instructor_id  
FROM   course c, section s  
WHERE  c.course_no = s.course_no
```

[Figure 7.6](#) shows a partial listing of the INSTRUCTOR table. The INSTRUCTOR_ID column is the primary key of the table and is the common column with the SECTION table. Every row in the SECTION table that has a value for the INSTRUCTOR_ID column must have one corresponding row in the INSTRUCTOR table. A particular INSTRUCTOR_ID in the INSTRUCTOR table may have zero, one, or multiple rows in the SECTION table.

To formulate the SQL statement, follow the same steps performed in [Lab 7.1](#). First, determine the columns and tables needed for output. Then confirm whether a one-to-one or a one-to-many relationship exists between the tables to accomplish the join. The changes to the previous SQL statement are indicated here in bold.

```
SELECT c.course_no, s.section_no,
       c.description, s.location,
       s.instructor_id, i.last_name,
       i.first_name
FROM course c, section s,
instructor i
WHERE c.course_no = s.course_no
AND s.instructor_id =
i.instructor_id
```

303

304

FIGURE 7.5 Result of join between COURSE and SECTION tables

COURSE_NO	SECTION_NO	DESCRIPTION	LOCATION	INSTRUCTOR_ID
10	2	Technology Concepts	L214	102
20	2	Intro to Information Systems	L210	103
20	4	Intro to Information Systems	L214	104
20	7	Intro to Information Systems	L509	105
20	8	Intro to Information Systems	L210	106
25	1	Intro to Programming	M311	107
25	2	Intro to Programming	L210	108
25	3	Intro to Programming	L507	101
25	4	Intro to Programming	L214	102
25	5	Intro to Programming	L509	103
25	6	Intro to Programming	L509	104
25	7	Intro to Programming	L210	105
25	8	Intro to Programming	L509	106
25	9	Intro to Programming	L507	107
100	1	Hands-On Windows	L214	102
100	2	Hands-On Windows	L500	103
100	3	Hands-On Windows	L210	104

FIGURE 7.6 The INSTRUCTOR table

INSTRUCTOR_ID	...	FIRST_NAME	LAST_NAME						
101	Mr	Fernand	Hanks
102	Mr	Tom	Wojick
103	Ms	Nina	Schorin
104	Mr	Gary	Pertez
105	Ms	Anita	Morris
106	Rev	Todd	Smythe
107	Dr	Marilyn	Frantzen
108	Mr	Charles	Lowry
109	Hon	Rick	Chow
110	Ms	Irene	Willig

The join yields the result shown in [Figure 7.7](#). The three-table join result now includes the instructor's first and last names. For example, INSTRUCTOR_ID 102 is listed multiple times in the SECTION table. This instructor teaches several sections; therefore, the INSTRUCTOR_ID's corresponding first and last names are repeated in the result.

304

FIGURE 7.7 Result of join between the COURSE, SECTION, and INSTRUCTOR tables

COU...	SEC...	DESCRIPTION	LOC...	INSTRUC...	LAST_NAME	FIRST_NAME
10	2	Technology Concepts	L214	102	Wojick	Tom
20	2	Intro to Information Systems	L210	103	Schorin	Nina
20	4	Intro to Information Systems	L214	104	Pertez	Gary
20	7	Intro to Information Systems	L509	105	Morris	Anita
20	8	Intro to Information Systems	L210	106	Smythe	Todd
25	1	Intro to Programming	M311	107	Frantzen	Marilyn
25	2	Intro to Programming	L210	108	Lowry	Charles
25	3	Intro to Programming	L507	101	Hanks	Fernand
25	4	Intro to Programming	L214	102	Wojick	Tom
25	5	Intro to Programming	L509	103	Schorin	Nina
25	6	Intro to Programming	L509	104	Pertez	Gary
25	7	Intro to Programming	L210	105	Morris	Anita
25	8	Intro to Programming	L509	106	Smythe	Todd
25	9	Intro to Programming	L507	107	Frantzen	Marilyn
100	1	Hands-On Windows	L214	102	Wojick	Tom
100	2	Hands-On Windows	L509	103	Schorin	Nina

ANSI Join Syntax for Joining Three or More Tables

A join across three tables can be expressed with the ANSI join syntax. Create the first join between the COURSE and SECTION tables via the JOIN keyword and the ON clause. To this result, you add the next table and join condition. The set of parentheses around the ON clause is optional.

```
SELECT c.course_no, s.section_no,
       c.description, s.location,
```

```
    s.instructor_id, i.last_name,  
i.first_name  
FROM course c JOIN section s  
ON (c.course_no = s.course_no)  
JOIN instructor i  
ON (s.instructor_id =  
i.instructor_id)
```

Alternatively, the query can be expressed with the USING clause. The table and column aliases in the SELECT and FROM clauses are optional, but the parentheses in the USING clause are required.

```
SELECT course_no, s.section_no,  
c.description, s.location,  
instructor_id, i.last_name,  
i.first_name  
FROM course c JOIN section s  
USING (course_no)  
JOIN instructor i  
USING (instructor_id)
```

Multicolumn Joins

The basic steps of a multicolumn join do not differ from the previous examples. The only variation is to make multicolumn keys part of the join criteria.

305

One of the multikey column examples in the schema is the GRADE table. The primary key of the table consists of the four columns STUDENT_ID, SECTION_ID, GRADE_CODE_OCCURRENCE, and GRADE_TYPE_CODE. The GRADE table also has two foreign keys: the GRADE_TYPE_CODE column, referencing the GRADE_TYPE_WEIGHT table, and the multi-column foreign keys STUDENT_ID and SECTION_ID, referencing the ENROLLMENT table.

To help you understand the data in the table, examine a set of sample records for a particular student. The student with ID 220 is enrolled in SECTION_ID 119 and has nine records in the GRADE table: four homework assignments (HM), two quizzes (QZ), one midterm (MT), one final examination (FI), and one participation (PA) grade.

```
SELECT student_id, section_id, grade_type_code type,
       grade_code_occurrence no,
       numeric_grade indiv_gr
FROM grade
WHERE student_id = 220
      AND section_id = 119
```

STUDENT_ID	SECTION_ID	TY	NO	INDIV_GR
220	119	FI	1	85
220	119	HM	1	84
220	119	HM	2	84
220	119	HM	3	74
220	119	HM	4	74
220	119	MT	1	88
220	119	PA	1	91
220	119	QZ	1	92
220	119	QZ	2	91

9 rows selected.

The next SQL query joins the GRADE table to the ENROLLMENT table, to include the values of the ENROLL_DATE column in the result set. All the changes from the previous SQL query are indicated in bold.

```
SELECT g.student_id, g.section_id,
       g.grade_type_code type,
       g.grade_code_occurrence no,
       g.numeric_grade indiv_gr
       ,TO_CHAR(e.enroll_date, 'MM/DD/
YY') enrolldt
FROM grade g, enrollment e
WHERE g.student_id = 220
AND g.section_id = 119
AND g.student_id = e.student_id
AND g.section_id = e.section_id
```

306

307

```
SELECT g.student_id, g.section_id,
       g.grade_type_code type,
       g.grade_code_occurrence no,
       g.numeric_grade indiv_gr,
       TO_CHAR(e.enroll_date, 'MM/DD/YY') enrolldt
FROM grade g, enrollment e
WHERE g.student_id = 220
AND g.section_id = 119
AND g.student_id = e.student_id
AND g.section_id = e.section_id
```

STUDENT_ID	SECTION_ID	TY	NO	INDIV_GR	ENROLLDT
220	119	FI	1	85	02/16/07
220	119	HM	1	84	02/16/07
220	119	HM	2	84	02/16/07
220	119	HM	3	74	02/16/07
220	119	HM	4	74	02/16/07
220	119	MT	1	88	02/16/07
220	119	PA	1	91	02/16/07
220	119	QE	1	92	02/16/07
220	119	QE	2	91	02/16/07

9 rows selected.

To join between the tables ENROLLMENT and GRADE, use both the SECTION_ID and STUDENT_ID columns. These two columns represent the primary key of the ENROLLMENT table and the foreign key of the GRADE table; thus a one-to-many relationship between the tables exists.

The values for the ENROLL_DATE column are repeated, because for each individual grade you have one row showing ENROLL_DATE in the ENROLLMENT table.

EXPRESSING MULTICOLUMN JOINS USING THE ANSI JOIN SYNTAX

A join involving multiple columns on a table requires the columns to be listed in the ON or the USING clause as join criteria. The next SQL statement shows the ON clause.

```
SELECT g.student_id, g.section_id,
       g.grade_type_code type,
       g.grade_code_occurrence no,
       g.numeric_grade indiv_gr,
       TO_CHAR(e.enroll_date, 'MM/DD/YY')
enrolldt
FROM grade g JOIN enrollment e
```

```
ON (g.student_id = e.student_id
AND g.section_id = e.section_id)
WHERE g.student_id = 220
AND g.section_id = 119
```

When you write the query with the USING clause, you list the join columns separated by commas.

```
SELECT student_id, section_id,
grade_type_code type,
grade_code_occurrence no,
numeric_grade indiv_gr,
TO_CHAR(enroll_date, 'MM/DD/YY')
enrolldt
FROM grade JOIN enrollment
USING (student_id, section_id)
WHERE student_id = 220
AND section_id = 119
```

307

308

Joining Across Many Tables

Joining across multiple tables involves repeating the same steps of a two-join or three-join table: You join the first two tables and then join the result to each subsequent table, using the common column(s). You repeat this until all the tables are joined.

To join n tables together, you need at least $n-1$ join conditions. For example, to join five tables, you need at

least four join conditions. If your join deals with tables containing multicolumn keys, you will obviously need to include these multiple columns as part of the join condition.



The Oracle optimizer determines the order in which the tables are joined based on the join condition, the indexes on the table, and the various statistics about the tables (such as the number of rows or the number of distinct values in each column). The join order can have a significant impact on the performance of multitable joins. You can learn more about this topic and how to influence the optimizer in [Chapter 18](#).

The ANSI Join Versus the Traditional Join Syntax

You might wonder which one of the join syntax options is better. The ANSI join syntax has a number of advantages.

- It helps you easily identify the join criteria and the filtering condition.

- You avoid an accidental Cartesian product because you must explicitly specify the join criteria, and any missing join conditions become evident because an error is generated.
- The syntax is easy to read and understand.
- The USING clause requires less typing, but the data types of the columns must match.
- SQL is understood by other ANSI-compliant non-Oracle databases.

Although the traditional join syntax, with the columns separated by commas in the FROM clause and the join condition listed in the WHERE clause, may become the old way of writing SQL, you must nevertheless familiarize yourself with this syntax because millions of SQL statements already use it, and it clearly performs its intended purpose.

The ANSI join syntax has some distinct functional advantages over the traditional join syntax when it comes to outer joins, which you can learn about in [Chapter 10](#).

Using SQL Developer's Query Builder

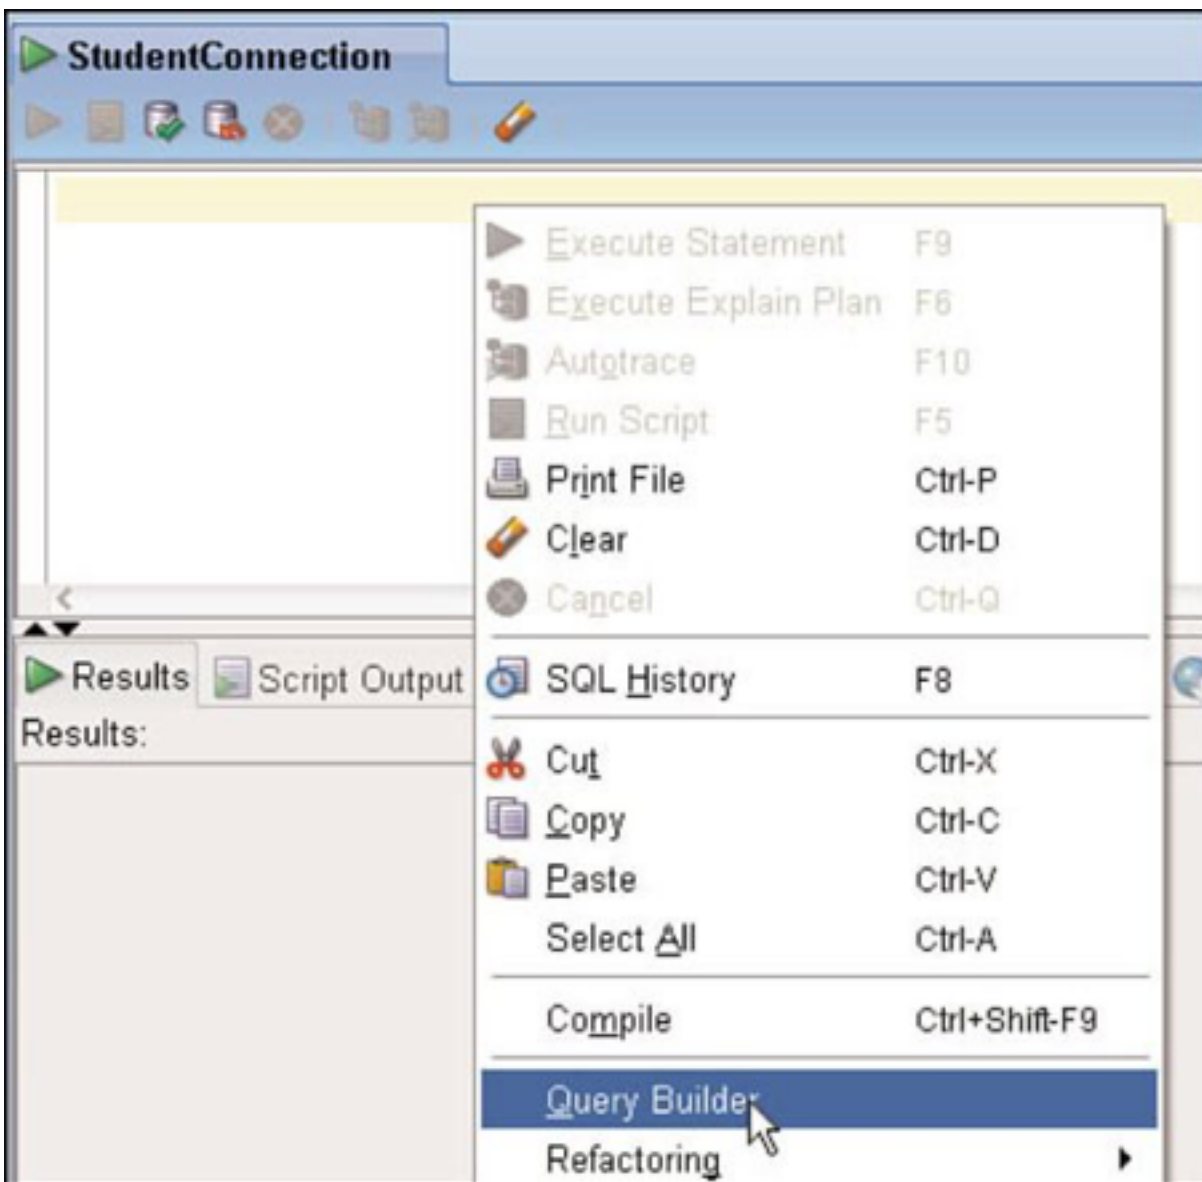
Instead of typing your SQL statement and figuring out the joining columns, SQL Developer performs these

tasks for you with the help of the Query Builder. As [Figure 7.8](#) shows, you invoke the Query Builder from the SQL Worksheet by right-clicking and selecting Query Builder.

308

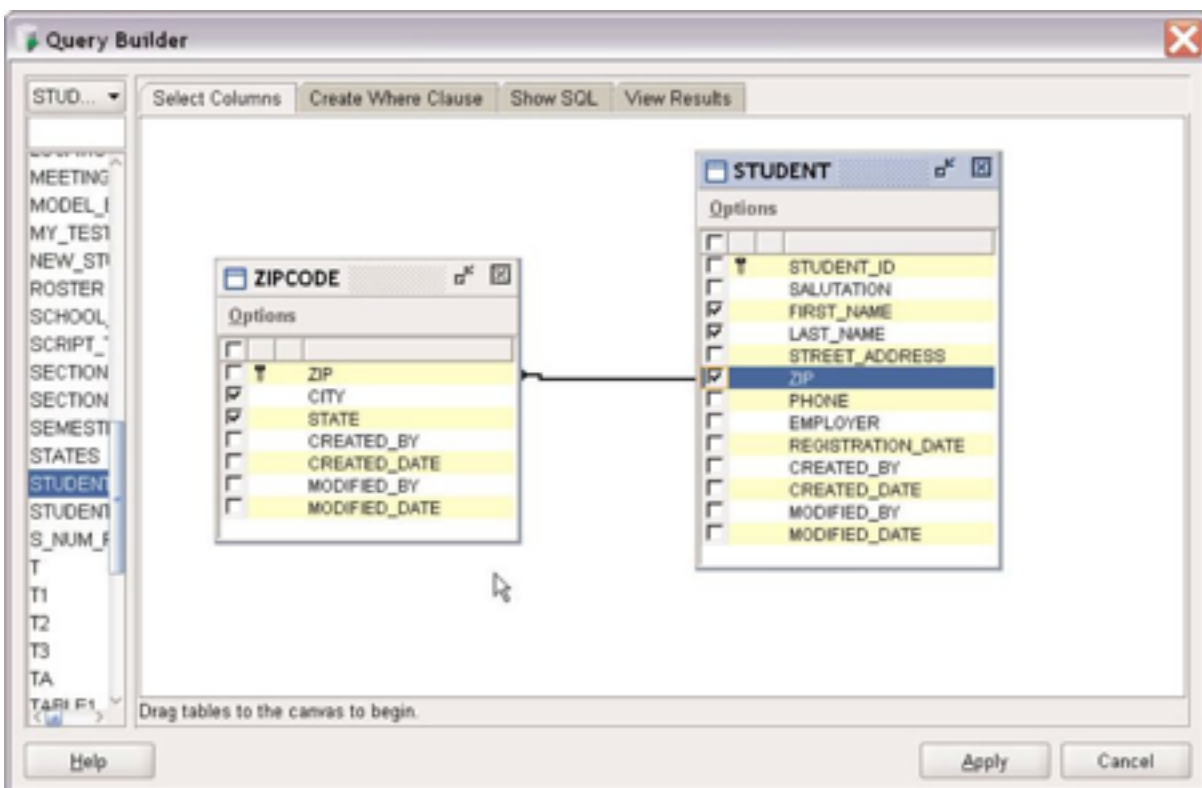
309

FIGURE 7.8 Query Builder



Query Builder allows you to drag the desired tables to the canvas. It builds the join criteria based on the existing primary/foreign key relationship and visually displays the relationship line between the tables (see [Figure 7.9](#)). You can check the columns you want to include in the query. To narrow down the query output, use the Create Where Clause tab at the top of the screen. SQL Developer generates a SQL statement you can view and execute.

FIGURE 7.9 Query Builder canvas



309

Different Types of Joins

Most of the joins you will come across are based on equality, with the equijoin being the most dominant. In this chapter, you have learned about equijoins; there are other types of joins you must become familiar with, most notably the self-join, the nonequijoin, and the outer join. [Table 7.1](#) lists the various types of joins.

TABLE 7.1 Types of Joins

JOIN TYPE	BASE OF JOIN CONDITION	LEARN ABOUT IT	SYNTAX
Equijoin or inner join	Equality	This chapter	Traditional comma-separated join or ANSI join syntax (including the optional INNER keyword).
Natural join	Equality	This chapter	NATURAL JOIN keyword.
Cross-join or Cartesian product	No join condition	This chapter	Traditional comma-separated with the missing join condition in the WHERE clause or CROSS JOIN keyword.
Self-join	Equality	Chapter 10	See equijoin or inner join.
Outer join (left, right, full)	Equality and extending the result set	Chapter 10	OUTER JOIN keywords or outer join operator (+).

Nonequijoin	Nonequality of values	Chapter 10	Traditional comma-separated join or ANSI join syntax with the ON clause. Join criteria is not based on equality.
-------------	--------------------------	--------------------------------	--

LAB 7.2 EXERCISES

- a) Display the student ID, course number, and section number of enrolled students where the instructor of the section lives in zip code 10025. In addition, the course should not have any prerequisites.
- b) Produce the mailing addresses for instructors who taught sections that started in June 2007.
- c) List the student IDs of enrolled students living in Connecticut.
- d) Show all the grades student Fred Crocitto received for SECTION_ID 86.
- e) List the final examination grades for all enrolled Connecticut students of course number 420. (Note that final examination grade does not mean final grade.)

- f) Display the LAST_NAME, STUDENT_ID, PERCENT_OF_FINAL_GRADE, GRADE_TYPE_CODE, and NUMERIC_GRADE columns for students who received 80 or less for their class project (GRADE_TYPE_CODE = 'PJ'). Order the result by student last name.

310

311

LAB 7.2 EXERCISE ANSWERS

- a) Display the student ID, course number, and section number of enrolled students where the instructor of the section lives in zip code 10025. In addition, the course should not have any prerequisites.

ANSWER: This query involves joining four tables. The course number is found in the SECTION and COURSE tables and the PREREQUISITE column in the COURSE table. To determine the zip code of an instructor, use the INSTRUCTOR table. To choose only enrolled students, join to the ENROLLMENT table.

```
SELECT c.course_no, s.section_no, e.student_id
FROM course c, section s, instructor i, enrollment e
WHERE c.prerequisite IS NULL
AND c.course_no = s.course_no
AND s.instructor_id = i.instructor_id
AND i.zip = '10025'
AND s.section_id = e.section_id
COURSE_NO SECTION_NO STUDENT_ID
-----
      10           2       128
     146           2       117
     146           2       140
...
      20           8       158
      20           8       199

12 rows selected.
```

To obtain this result, build the four-table join as you would any other join, step by step. First, start with one of the tables, such as the COURSE table.

```
SELECT course_no
FROM course
WHERE prerequisite IS NULL
```

For each of these courses, find the corresponding sections when you join the COURSE table with the SECTION table. Notice the bolded additions to the SQL statement.

```
SELECT c.course_no, s.section_no
FROM course c, section s
```

```
WHERE c.prerequisite IS NULL
AND c.course_no = s.course_no
```

Then include instructors who live in zip code 10025. The common column between SECTION and INSTRUCTOR is INSTRUCTOR_ID.

```
SELECT c.course_no, s.section_no
FROM course c, section s,
instructor i
WHERE c.prerequisite IS NULL
AND c.course_no = s.course_no
AND s.instructor_id =
i.instructor_id
AND i.zip = '10025'
```

311

312

Finally, join the results of the ENROLLMENT table via the SECTION_ID column, which leads you to the solution shown previously.

Instead of using the traditional join syntax to obtain the result, you can use the ANSI join syntax. The query may look similar to the following statement.

```
SELECT course_no, section_no,
student_id
FROM course JOIN section
USING (course_no)
```

```
JOIN instructor
USING (instructor_id)
JOIN enrollment
USING (section_id)
WHERE prerequisite IS NULL
AND zip = '10025'
```

The following is another possible alternative using the ANSI join syntax, using the ON condition instead of the USING clause.

```
SELECT c.course_no,
s.section_no, e.student_id
FROM course c JOIN section s
ON (c.course_no = s.course_no)
JOIN instructor i
ON (s.instructor_id =
i.instructor_id)
JOIN enrollment e
ON (s.section_id = e.section_id)
WHERE c.prerequisite IS NULL
AND i.zip = '10025'
```

- b)** Produce the mailing addresses for instructors who taught sections that started in June 2007.

ANSWER: This solution requires the join of three tables: You join the INSTRUCTOR,

SECTION, and ZIPCODE tables to produce the mailing list.

```

SELECT i.first_name || ' ' || i.last_name name,
       i.street_address, z.city || ', ' || z.state
       || ' ' || i.zip "City State Zip",
       TO_CHAR(s.start_date_time, 'MM/DD/YY') start_dt,
       section_id sect
FROM instructor i, section s, zipcode z
WHERE i.instructor_id = s.instructor_id
      AND i.zip = z.zip
      AND s.start_date_time >=
          TO_DATE('01-JUN-2007', 'DD-MON-YYYY')
      AND s.start_date_time <
          TO_DATE('01-JUL-2007', 'DD-MON-YYYY')

```

NAME	STREET_ADDRESS	City	State	Zip	START_DT	SECT
Fernand Hanks	100 East 87th	New York, NY	10015	06/02/07	117	
Anita Morris	34 Maiden Lane	New York, NY	10015	06/11/07	83	
Anita Morris	34 Maiden Lane	New York, NY	10015	06/12/07	91	
Anita Morris	34 Maiden Lane	New York, NY	10015	06/02/07	113	
...						
Gary Pertez	34 Sixth Ave	New York, NY	10035	06/12/07	90	
Gary Pertez	34 Sixth Ave	New York, NY	10035	06/10/07	120	
Gary Pertez	34 Sixth Ave	New York, NY	10035	06/03/07	143	
Gary Pertez	34 Sixth Ave	New York, NY	10035	06/12/07	151	

312

313

One of the first steps in solving this problem is to determine what columns and tables are involved. Look at the schema diagram in [Appendix D](#), “STUDENT Database Schema,” or refer to the table and column comments listed in [Appendix E](#), “Table and Column Descriptions.”

In this example, the instructor’s last name, first name, street address, and zip code are found in the INSTRUCTOR table. CITY, STATE, and ZIP

are columns in the ZIPCODE table. The join also needs to include the SECTION table because the column START_DATE_TIME lists the date and time when the individual sections started. The next step is to determine the common columns. The ZIP column is the common column between the INSTRUCTOR and ZIPCODE tables. For every value in the ZIP column of the INSTRUCTOR table you have one corresponding ZIP value in the ZIPCODE table. For every value in the ZIPCODE table there may be zero, one, or multiple records in the INSTRUCTOR table. The join returns only the matching records.

The other common column is INSTRUCTOR_ID in the SECTION and INSTRUCTOR tables. Only instructors who teach have one or more rows in the SECTION table. Any section that does not have an instructor assigned is not taught.

As always, the query can be expressed with one of the ANSI join syntax variations.

```
SELECT first_name || ' ' ||  
last_name name,
```

```
street_address, city || ', ' ||
state
|| ' ' || zip "City State Zip",
TO_CHAR(start_date_time, 'MM/DD/
YY') start_dt,
section_id sect
FROM instructor JOIN section s
USING (instructor_id)
JOIN zipcode
USING (zip)
WHERE start_date_time
>=TO_DATE('01-JUN-2007', 'DD-MON-
YYYY')
AND start_date_time <
TO_DATE('01-JUL-2007', 'DD-MON-
YYYY')
```

You see that some instructors are teaching multiple sections. To see only the distinct addresses, use the **DISTINCT** keyword and drop the **START_DATE_TIME** and **SECTION_ID** columns from the **SELECT** list.

- c) List the student IDs of enrolled students living in Connecticut.

ANSWER: Only students enrolled in classes are in the result; any student who does not have a

row in the ENROLLMENT table is not considered enrolled. STUDENT_ID is the common column between the STUDENT and ENROLLMENT tables. The STATE column is in the ZIPCODE table. The common column between the STUDENT and ZIPCODE tables is the ZIP column.

```
SELECT student_id
FROM student JOIN enrollment
USING (student_id)
JOIN zipcode
USING (zip)
WHERE state = 'CT'
```

313

314

STUDENT_ID

220

270

270

...

210

154

13 rows selected.

Because students can be enrolled in more than one class, add the DISTINCT keyword if you want to display each STUDENT_ID once.

Following is the SQL statement, expressed using the traditional join syntax. SELECT s.student_id

```
FROM student s, enrollment e,  
zipcode z
```

```
WHERE s.student_id =  
e.student_id
```

```
AND s.zip = z.zip
```

```
AND z.state = 'CT'
```

- d) Show all the grades student Fred Crocitto received for SECTION_ID 86.

ANSWER: The grades for each section and student are stored in the GRADE table. The primary key of the GRADE table consists of the STUDENT_ID, SECTION_ID, GRADE_TYPE_CODE, and GRADE_CODE_OCCURRENCE columns. This means a student, such as Fred Crocitto, has multiple grades for each grade type.

```
SELECT s.first_name || ' ' ||  
s.last_name name,  
e.section_id, g.grade_type_code,  
g.numeric_grade grade  
FROM student s JOIN enrollment e  
ON (s.student_id = e.student_id)
```

```

JOIN grade g
ON (e.student_id = g.student_id
AND e.section_id = g.section_id)
WHERE s.last_name = 'Crocitto'
AND s.first_name = 'Fred'
AND e.section_id = 86

```

The SQL statement using the traditional join syntax may look similar to the following query.

```

SELECT s.first_name|| ' ' ||
s.last_name name,
e.section_id, g.grade_type_code,
g.numeric_grade grade
FROM student s, enrollment e,
grade g
WHERE s.last_name = 'Crocitto'
AND s.first_name = 'Fred'
AND e.section_id = 86
AND s.student_id = e.student_id
AND e.student_id = g.student_id
AND e.section_id = g.section_id

```

314

```

SELECT s.first_name|| ' ' || s.last_name name,
e.section_id, g.grade_type_code,
g.numeric_grade grade
FROM student s, enrollment e, grade g
WHERE s.last_name = 'Crocitto'
AND s.first_name = 'Fred'
AND e.section_id = 86
AND s.student_id = e.student_id
AND e.student_id = g.student_id
AND e.section_id = g.section_id

NAME                SECTION_ID OR GRADE
-----
Fred Crocitto       86 F2      85
...
Fred Crocitto       86 Q5      90
Fred Crocitto       86 Q5      84
Fred Crocitto       86 Q5      97
Fred Crocitto       86 Q5      97

11 rows selected.

```

315

To build up the SQL statement step by step, you might want to start with the STUDENT table and select the record for Fred Crocitto.

```
SELECT last_name, first_name
FROM student
WHERE last_name = 'Crocitto'
AND first_name = 'Fred'
```

Next, choose the section with ID 86, in which Fred is enrolled. The common column between the two tables is STUDENT_ID.

```
SELECT s.first_name||' '||
s.last_name name
, e.section_id
FROM student s, enrollment e
WHERE s.last_name = 'Crocitto'
AND s.first_name = 'Fred'
AND e.section_id = 86
AND s.student_id = e.student_id
```

Next, retrieve the individual grades from the GRADE table. The common columns between the GRADE table and the ENROLLMENT table are SECTION_ID and STUDENT_ID. They represent the primary key in the ENROLLMENT

table and are foreign keys in the GRADE table. Both columns need to be in the WHERE clause.

If you want to expand the query, add the DESCRIPTION column of the GRADE_TYPE table for each GRADE_TYPE_CODE. The common column between the tables GRADE and GRADE_TYPE is GRADE_TYPE_CODE.

```
SELECT s.first_name||' '||
s.last_name name,
e.section_id, g.grade_type_code
grade,
g.numeric_grade, gt.description
FROM student s, enrollment e,
grade g, grade_type gt
WHERE s.last_name = 'Crocitto'
AND s.first_name = 'Fred'
AND e.section_id = 86
AND s.student_id = e.student_id
AND e.student_id = g.student_id
AND e.section_id = g.section_id
AND g.grade_type_code =
gt.grade_type_code
```

If you also show the COURSE_NO column, join to the SECTION table via the ENROLLMENT table column SECTION_ID.

```

SELECT s.first_name||' '|| s.last_name name,
       e.section_id, g.grade_type_code,
       g.numeric_grade grade, gt.description,
       sec.course_no
FROM student s, enrollment e, grade g, grade_type gt,
       section sec
WHERE s.last_name = 'Crocitto'
      AND s.first_name = 'Fred'
      AND e.section_id = 86
      AND s.student_id = e.student_id
      AND e.student_id = g.student_id
      AND e.section_id = g.section_id
      AND g.grade_type_code = gt.grade_type_code
      AND e.section_id = sec.section_id
NAME                SECTION_ID GR GRADE DESCRIPTION COURSE_NO
-----
Fred Crocitto       86 FI      85 Final          25
...
Fred Crocitto       86 QZ      90 Quiz           25
Fred Crocitto       86 QZ      84 Quiz           25
Fred Crocitto       86 QZ      97 Quiz           25
Fred Crocitto       86 QZ      97 Quiz           25

11 rows selected.

```

- e) List the final examination grades for all enrolled Connecticut students of course number 420.
(Note that final examination grade does not mean final grade.)

ANSWER: To find the answer, you need to join five tables. The needed joins are the ZIPCODE table with the STUDENT table to determine the Connecticut students and the STUDENT and ENROLLMENT tables to determine the SECTION_IDS in which the students are

enrolled. From these SECTION_IDs, you include only sections where the course number equals 420. This requires a join of the ENROLLMENT table to the SECTION table. Finally, join the ENROLLMENT table to the GRADE table to display the grades.

```
SELECT e.student_id, sec.course_no, g.numeric_grade
FROM student stud, zipcode z,
      enrollment e, section sec, grade g
WHERE stud.zip = z.zip
      AND z.state = 'CT'
      AND stud.student_id = e.student_id
      AND e.section_id = sec.section_id
      AND e.section_id = g.section_id
      AND e.student_id = g.student_id
      AND sec.course_no = 420
      AND g.grade_type_code = 'FI'
STUDENT_ID COURSE_NO NUMERIC_GRADE
-----
          196          420          84
          198          420          85

2 rows selected.
```

316

317

You can list any of the columns you find relevant to solving the query. For this solution, the columns STUDENT_ID, COURSE_NO, and NUMERIC_GRADE were chosen.

The query can also be expressed with the ANSI join syntax, as in the following example, which shows the USING clause.

```
SELECT student_id, course_no,  
numeric_grade  
FROM student JOIN zipcode  
USING (zip)  
JOIN enrollment  
USING (student_id)  
JOIN section  
USING (section_id)  
JOIN grade g  
USING (section_id, student_id)  
WHERE course_no = 420  
AND grade_type_code = 'FI'  
AND state = 'CT'
```

- f)** Display the LAST_NAME, STUDENT_ID, PERCENT_OF_FINAL_GRADE, GRADE_TYPE_CODE, and NUMERIC_GRADE columns for students who received 80 or less for their class project (GRADE_TYPE_CODE = 'PJ'). Order the result by student last name.

ANSWER: Join the tables GRADE_TYPE_WEIGHT, GRADE, ENROLLMENT, and STUDENT.

The column `PERCENT_OF_FINAL_GRADE` of the `GRADE_TYPE_WEIGHT` table stores the weighted percentage a particular grade has on the final grade. One of the foreign keys of the `GRADE` table is the combination of `GRADE_TYPE_CODE` and `SECTION_ID`; these columns represent the primary key of the `GRADE_TYPE_WEIGHT` table.

To include the student's last name, you have two choices. Either follow the primary and foreign key relationships by joining the tables `GRADE` and `ENROLLMENT` via the `STUDENT_ID` and `SECTION_ID` columns and then join the `ENROLLMENT` table to the `STUDENT` table via the `STUDENT_ID` column, or skip the `ENROLLMENT` table and join `GRADE` directly to the `STUDENT` table via `STUDENT_ID`. Examine the first option of joining to the `ENROLLMENT` table and then joining it to the `STUDENT` table.

```
SELECT g.student_id,  
       g.section_id,  
       gw.percent_of_final_grade pct,  
       g.grade_type_code,
```

```

    g.numeric_grade grade,
s.last_name
FROM grade_type_weight gw, grade
g,
enrollment e, student s
WHERE g.grade_type_code = 'PJ'
AND gw.grade_type_code =
g.grade_type_code
AND gw.section_id = g.section_id
AND g.numeric_grade <= 80
AND g.section_id = e.section_id
AND g.student_id = e.student_id
AND e.student_id = s.student_id
ORDER BY s.last_name

```

317

```

SELECT g.student_id, g.section_id,
      gw.percent_of_final_grade pct, g.grade_type_code,
      g.numeric_grade grade, s.last_name
FROM grade_type_weight gw, grade g,
enrollment e, student s
WHERE g.grade_type_code = 'PJ'
AND gw.grade_type_code = g.grade_type_code
AND gw.section_id = g.section_id
AND g.numeric_grade <= 80
AND g.section_id = e.section_id
AND g.student_id = e.student_id
AND e.student_id = s.student_id
ORDER BY s.last_name

```

318

STUDENT_ID	SECTION_ID	PCT GR	GRADE LAST_NAME
245	82	75 PJ	77 Dalvi
176	115	75 PJ	76 Satterfield
244	82	75 PJ	76 Wilson
248	155	75 PJ	76 Zapulla

4 rows selected.

SKIPPING THE PRIMARY/FOREIGN KEY PATH

The second choice is to join the STUDENT_ID from the GRADE table directly to the STUDENT_ID of the STUDENT table, thus skipping the ENROLLMENT table entirely. The following query returns the same result as the query in the preceding section.

```
SELECT g.student_id, g.section_id,
       gw.percent_of_final_grade pct, g.grade_type_code,
       g.numeric_grade grade, s.last_name
FROM grade_type_weight gw, grade g,
     student s
WHERE g.grade_type_code = 'PJ'
      AND gw.grade_type_code = g.grade_type_code
      AND gw.section_id = g.section_id
      AND g.numeric_grade <= 80
      AND g.student_id = s.student_id
ORDER BY s.last_name
```

STUDENT_ID	SECTION_ID	PCT	GR	GRADE	LAST_NAME
245	82	75	PJ	77	Dalvi
176	115	75	PJ	76	Satterfield
244	82	75	PJ	76	Wilson
248	155	75	PJ	76	Zapulla

4 rows selected.

This shortcut is perfectly acceptable, even if it does not follow the primary key/foreign key relationship path. In this case, you can be sure not to build a Cartesian

product because you can guarantee only one STUDENT_ID in the STUDENT table for every STUDENT_ID in the GRADE table. In addition, it also eliminates a join; thus, the query executes a little faster and requires fewer resources. The effect is probably fairly negligible with this small result set.

318

319

Lab 7.2 Quiz

In order to test your progress, you should be able to answer the following questions.

1) Which SQL statement shows the sections that have instructors assigned to them?

_____ **a)** `SELECT c.course_no,
s.section_id,
i.instructor_id

FROM course c, section s,
instructor i
WHERE c.course_no =
s.course_no
AND i.instructor_id =
s.section_id`

_____ **b)** `SELECT c.course_no,
s.section_id,
i.instructor_id`

```
FROM course c, section s,  
instructor i  
WHERE c.course_no =  
s.course_no  
AND i.instructor_id =  
s.instructor_id
```

_____ **c)** SELECT course_no,
section_id,
instructor.instructor_id

FROM section, instructor
WHERE
instructor.instructor_id =
section.section_id

_____ **d)** SELECT c.section_id,
i.instructor_id

FROM course c, instructor
i
WHERE i.instructor_id =
c.section_id

_____ **e)** SELECT c.course_no,
i.instructor_id

FROM course c JOIN
instructor
USING (instructor_id)

2) How do you resolve the Oracle error ORA-00918: column ambiguously defined?

- a) _____ Correct the join criteria and WHERE clause condition.
- b) _____ Choose another column.
- c) _____ Add the correct table alias.
- d) _____ Correct the spelling of the column name.

3) Joins involving multiple columns must always follow the primary key/foreign key relationship path.

- a) _____ True
- b) _____ False

4) Find the error(s) in the following SQL statement.

```
1 SELECT g.student_id, s.section_id,  
2    g.numeric_grade, s.last_name  
3 FROM grade g,  
4    enrollment e, student s  
5 WHERE g.section_id = e.section_id
```

6 AND g.student_id = e.student_id

7 AND s.student_id = e.student_id

8 AND s.student_id = 248

9 AND e.section_id = 155

319

320

a) _____ Line 1

b) _____ Line 5

c) _____ Line 6

d) _____ Lines 5 and 6

e) _____ Lines 1, 5, and 6

f) _____ No error

5) Equijoins are the most common type of joins and are always based on equality of values.

a) _____ True

b) _____ False

6) To join four tables, you must have at least three join conditions.

a) _____ True

b) _____ False

320

WORKSHOP

The projects in this section are meant to prompt you to utilize all the skills you have acquired throughout this chapter. The answers to these projects can be found at the companion Web site to this book, located at www.oraclesqlbyexample.com.

- 1) Select the course description, section number, and location for sections meeting in location L211.
- 2) Show the course description, section number, and starting date and time of the courses Joseph German is taking.
- 3) List the instructor ID, last name of the instructor, and section ID of sections where class participation contributes to 25 percent of the total grade. Order the result by the instructor's last name.
- 4) Display the first and last names of students who received 99 or more points on the class project.
- 5) Select the grades for quizzes of students living in zip code 10956.

- 6) List the course number, section number, and instructor first and last names for classes with course number 350 as a prerequisite.
- 7) What problem do the following two SELECT statements solve? Explain the difference between the two results.

```
SELECT stud.student_id,  
i.instructor_id,  
stud.zip, i.zip  
FROM student stud, instructor i  
WHERE stud.zip = i.zip  
SELECT stud.student_id,  
i.instructor_id,  
stud.zip, i.zip  
FROM student stud, enrollment e,  
section sec,  
instructor i  
WHERE stud.student_id =  
e.student_id  
AND e.section_id =  
sec.section_id  
AND sec.instructor_id =  
i.instructor_id  
AND stud.zip = i.zip
```